

# Running Airflow in Docker

This quick-start guide will allow you to quickly get Airflow up and running with [CeleryExecutor](#) in Docker.

! Caution

This procedure can be useful for learning and exploration. However, adapting it for use in real-world situations can be complicated. Making changes to this procedure will require specialized expertise in Docker & Docker Compose, and the Airflow community may not be able to help you.

For that reason, we recommend using Kubernetes with the [Official Airflow Community Helm Chart](#) when you are ready to run Airflow in production.

## Before you begin

This procedure assumes familiarity with Docker and Docker Compose. If you haven't worked with these tools before, you should take a moment to run through the [Docker Quick Start](#) (especially the section on [Docker Compose](#)) so you are familiar with how they work.

Follow these steps to install the necessary tools, if you have not already done so.

1. Install [Docker Community Edition \(CE\)](#) on your workstation. Depending on your OS, you may need to configure Docker to use at least 4.00 GB of memory for the Airflow containers to run properly. Please refer to the Resources section in the [Docker for Windows](#) or [Docker for Mac](#) documentation for more information.
2. Install [Docker Compose](#) v2.14.0 or newer on your workstation.

Older versions of `docker-compose` do not support all the features required by the Airflow `docker-compose.yaml` file, so double check that your version meets the minimum version requirements.

! Tip

The default amount of memory available for Docker on macOS is often not enough to get Airflow up and running. If enough memory is not allocated, it might lead to the webserver continuously restarting. You should allocate at least 4GB memory for the Docker Engine (ideally 8GB).

You can check if you have enough memory by running this command:

```
docker run --rm "debian:bullseye-slim" bash -c 'numfmt --to iec $(echo $(( $(getconf _PHYS_PAGES) * $(getconf PAGE_SIZE))))'
```

! Warning

### Running Airflow in Docker

[Before you begin](#)

[Fetching docker-compose.yaml](#)

[Initializing Environment](#)

[Setting the right Airflow user](#)

[Initialize the database](#)

[Cleaning-up the environment](#)

[Running Airflow](#)

[Accessing the environment](#)

[Running the CLI commands](#)

[Accessing the web interface](#)

[Sending requests to the REST API](#)

[Cleaning up](#)

[Using custom images](#)

[Special case - adding dependencies via requirements.txt file](#)

[Networking](#)

[FAQ: Frequently asked questions](#)

`ModuleNotFoundError: No module named 'XYZ'`

[What's Next?](#)

[Environment variables supported by Docker Compose](#)

Some operating systems (Fedora, ArchLinux, RHEL, Rocky) have recently introduced Kernel changes that result in Airflow in Docker Compose consuming 100% memory when run inside the community Docker implementation maintained by the OS teams.

This is an issue with backwards-incompatible containerd configuration that some of Airflow dependencies have problems with and is tracked in a few issues:

- [Moby issue](#)
- [Containerd issue](#)

There is no solution yet from the containerd team, but seems that installing [Docker Desktop on Linux](#) solves the problem as stated in [This comment](#) and allows to run Breeze with no problems.

## Fetching `docker-compose.yaml`

To deploy Airflow on Docker Compose, you should fetch [docker-compose.yaml](#).

```
curl -Lf0 'https://airflow.apache.org/docs/apache-airflow/2.8.1/docker-  
compose.yaml'
```

### ! Important

From July 2023 Compose V1 stopped receiving updates. We strongly advise upgrading to a newer version of Docker Compose, supplied `docker-compose.yaml` may not function accurately within Compose V1.

This file contains several service definitions:

- `airflow-scheduler` - The [scheduler](#) monitors all tasks and DAGs, then triggers the task instances once their dependencies are complete.
- `airflow-webserver` - The webserver is available at `http://localhost:8080`.
- `airflow-worker` - The worker that executes the tasks given by the scheduler.
- `airflow-triggerer` - The triggerer runs an event loop for deferrable tasks.
- `airflow-init` - The initialization service.
- `postgres` - The database.
- `redis` - [The redis](#) - broker that forwards messages from scheduler to worker.

Optionally, you can enable flower by adding `--profile flower` option, e.g. `docker compose --profile flower up`, or by explicitly specifying it on the command line e.g. `docker compose up flower`.

- `flower` - [The flower app](#) for monitoring the environment. It is available at `http://localhost:5555`.

All these services allow you to run Airflow with [CeleryExecutor](#). For more information, see [Architecture Overview](#).

Some directories in the container are mounted, which means that their contents are synchronized between your computer and the container.

- `./dags` - you can put your DAG files here.
- `./logs` - contains logs from task execution and scheduler.
- `./config` - you can add custom log parser or add `airflow_local_settings.py` to configure

cluster policy.

- `./plugins` - you can put your [custom plugins](#) here.

This file uses the latest Airflow image ([apache/airflow](#)). If you need to install a new Python library or system library, you can [build your image](#).

## Initializing Environment

Before starting Airflow for the first time, you need to prepare your environment, i.e. create the necessary files, directories and initialize the database.

### Setting the right Airflow user

On **Linux**, the quick-start needs to know your host user id and needs to have group id set to `0`. Otherwise the files created in `dags`, `logs` and `plugins` will be created with `root` user ownership. You have to make sure to configure them for the docker-compose:

```
mkdir -p ./dags ./logs ./plugins ./config
echo -e "AIRFLOW_UID=$(id -u)" > .env
```

See [Docker Compose environment variables](#)

For other operating systems, you may get a warning that `AIRFLOW_UID` is not set, but you can safely ignore it. You can also manually create an `.env` file in the same folder as `docker-compose.yaml` with this content to get rid of the warning:

```
AIRFLOW_UID=50000
```

### Initialize the database

On **all operating systems**, you need to run database migrations and create the first user account. To do this, run.

```
docker compose up airflow-init
```

After initialization is complete, you should see a message like this:

```
airflow-init_1      | Upgrades done
airflow-init_1      | Admin user airflow created
airflow-init_1      | 2.8.1
start_airflow-init_1 exited with code 0
```

The account created has the login `airflow` and the password `airflow` .

## Cleaning-up the environment

The docker-compose environment we have prepared is a “quick-start” one. It was not designed to be used in production and it has a number of caveats - one of them being that the best way to recover from any problem is to clean it up and restart from scratch.

The best way to do this is to:

- Run `docker compose down --volumes --remove-orphans` command in the directory you downloaded the `docker-compose.yaml` file
- Remove the entire directory where you downloaded the `docker-compose.yaml` file `rm -rf '<DIRECTORY>'`
- Run through this guide from the very beginning, starting by re-downloading the `docker-compose.yaml` file

## Running Airflow

Now you can start all services:

```
docker compose up
```

! Note

docker-compose is old syntax. Please check [Stackoverflow](#).

In a second terminal you can check the condition of the containers and make sure that no containers are in an unhealthy condition:

```
$ docker ps
CONTAINER ID   IMAGE                PORTS              COMMAND              CREATED
STATUS        PORTS             NAMES
247ebe6cf87a   apache/airflow:2.8.1 8080/tcp           "/usr/bin/dumb-init ..." 3 minutes ago
Up 3 minutes (healthy)           compose_airflow-worker_1
ed9b09fc84b1   apache/airflow:2.8.1 8080/tcp           "/usr/bin/dumb-init ..." 3 minutes ago
Up 3 minutes (healthy)           compose_airflow-scheduler_1
7cb1fb603a98   apache/airflow:2.8.1 0.0.0.0:8080->8080/tcp "/usr/bin/dumb-init ..." 3 minutes ago
Up 3 minutes (healthy)           compose_airflow-webserver_1
74f3bbe506eb   postgres:13          5432/tcp           "docker-entrypoint.s..." 18 minutes ago
Up 17 minutes (healthy)           compose_postgres_1
0bd6576d23cb   redis:latest          6379/tcp           "docker-entrypoint.s..." 10 hours ago
Up 17 minutes (healthy)           compose_redis_1
```

## Accessing the environment

After starting Airflow, you can interact with it in 3 ways:

- by running [CLI commands](#).
- via a browser using [the web interface](#).
- using [the REST API](#).

## Running the CLI commands

You can also run [CLI commands](#), but you have to do it in one of the defined `airflow-*` services. For example, to run `airflow info`, run the following command:

```
docker compose run airflow-worker airflow info
```

If you have Linux or Mac OS, you can make your work easier and download a optional wrapper scripts that will allow you to run commands with a simpler command.

```
curl -Lf0 'https://airflow.apache.org/docs/apache-airflow/2.8.1/airflow.sh'  
chmod +x airflow.sh
```

Now you can run commands easier.

```
./airflow.sh info
```

You can also use `bash` as parameter to enter interactive bash shell in the container or `python` to enter python container.

```
./airflow.sh bash
```

```
./airflow.sh python
```

## Accessing the web interface

Once the cluster has started up, you can log in to the web interface and begin experimenting with DAGs.

The webserver is available at: `http://localhost:8080`. The default account has the login `airflow` and the password `airflow`.

## Sending requests to the REST API

[Basic username password authentication](#) is currently supported for the REST API, which means you can use common tools to send requests to the API.

The webserver is available at: `http://localhost:8080` . The default account has the login `airflow` and the password `airflow` .

Here is a sample `curl` command, which sends a request to retrieve a pool list:

```
ENDPOINT_URL="http://localhost:8080/"
curl -X GET \
  --user "airflow:airflow" \
  "${ENDPOINT_URL}/api/v1/pools"
```

## Cleaning up

To stop and delete containers, delete volumes with database data and download images, run:

```
docker compose down --volumes --rmi all
```

## Using custom images

When you want to run Airflow locally, you might want to use an extended image, containing some additional dependencies - for example you might add new python packages, or upgrade airflow providers to a later version. This can be done very easily by specifying `build: .` in your `docker-compose.yaml` and placing a custom Dockerfile alongside your `docker-compose.yaml` . Then you can use `docker compose build` command to build your image (you need to do it only once). You can also add the `--build` flag to your `docker compose` commands to rebuild the images on-the-fly when you run other `docker compose` commands.

Examples of how you can extend the image with custom providers, python packages, apt packages and more can be found in [Building the image](#).

### ! Note

Creating custom images means that you need to maintain also a level of automation as you need to re-create the images when either the packages you want to install or Airflow is upgraded. Please do not forget about keeping these scripts. Also keep in mind, that in cases when you run pure Python tasks, you can use the [Python Virtualenv functions](#) which will dynamically source and install python dependencies during runtime. With Airflow 2.8.0 Virtualenvs can also be cached.

## Special case - adding dependencies via requirements.txt file

Usual case for custom images, is when you want to add a set of requirements to it - usually stored in `requirements.txt` file. For development, you might be tempted to add it dynamically when you are starting the original airflow image, but this has a number of side effects (for example your containers

will start much slower - each additional dependency will further delay your containers start up time). Also it is completely unnecessary, because docker compose has the development workflow built-in. You can - following the previous chapter, automatically build and use your custom image when you iterate with docker compose locally. Specifically when you want to add your own requirement file, you should do those steps:

1. Comment out the `image: ...` line and remove comment from the `build: .` line in the `docker-compose.yaml` file. The relevant part of the docker-compose file of yours should look similar to (use correct image tag):

```
#image: ${AIRFLOW_IMAGE_NAME:-apache/airflow:2.6.1}
build: .
```

2. Create `Dockerfile` in the same folder your `docker-compose.yaml` file is with content similar to:

```
FROM apache/airflow:2.6.1
ADD requirements.txt .
RUN pip install apache-airflow==${AIRFLOW_VERSION} -r requirements.txt
```

It is the best practice to install apache-airflow in the same version as the one that comes from the original image. This way you can be sure that `pip` will not try to downgrade or upgrade apache airflow while installing other requirements, which might happen in case you try to add a dependency that conflicts with the version of apache-airflow that you are using.

3. Place `requirements.txt` file in the same directory.

Run `docker compose build` to build the image, or add `--build` flag to `docker compose up` or `docker compose run` commands to build the image automatically as needed.

## Networking

In general, if you want to use Airflow locally, your DAGs may try to connect to servers which are running on the host. In order to achieve that, an extra configuration must be added in `docker-compose.yaml`. For example, on Linux the configuration must be in the section `services: airflow-worker` adding `extra_hosts: - "host.docker.internal:host-gateway"`; and use `host.docker.internal` instead of `localhost`. This configuration vary in different platforms. Please check the Docker documentation for [Windows](#) and [Mac](#) for further information.

## FAQ: Frequently asked questions

`ModuleNotFoundError: No module named 'XYZ'`

The Docker Compose file uses the latest Airflow image ([apache/airflow](#)). If you need to install a new Python library or system library, you can [customize and extend it](#).



## What’s Next?

From this point, you can head to the [Tutorials](#) section for further examples or the [How-to Guides](#) section if you’re ready to get your hands dirty.

## Environment variables supported by Docker Compose

Do not confuse the variable names here with the build arguments set when image is built. The `AIRFLOW_UID` build arg defaults to `50000` when the image is built, so it is “baked” into the image. On the other hand, the environment variables below can be set when the container is running, using `-` for example - result of `id -u` command, which allows to use the dynamic host runtime user id which is unknown at the time of building the image.

Variable	Description
<code>AIRFLOW_IMAGE_NAME</code>	Airflow Image to use.
<code>AIRFLOW_UID</code>	UID of the user to run Airflow containers as. Override if you want to

! Note

Before Airflow 2.2, the Docker Compose also had `AIRFLOW_GID` parameter, but it did not provide any additional functionality - only added confusion - so it has been removed.

Those additional variables are useful in case you are trying out/testing Airflow installation via Docker Compose. They are not intended to be used in production, but they make the environment faster to bootstrap for first time users with the most common customizations.

Variable	Description
<code>_AIRFLOW_WWW_USER_USERNAME</code>	Username for the administrator UI account. If this value is not set, the default is 'admin'.
<code>_AIRFLOW_WWW_USER_PASSWORD</code>	Password for the administrator UI account. Only used if <code>_AIRFLOW_WWW_USER_USERNAME</code> is set.
<code>_PIP_ADDITIONAL_REQUIREMENTS</code>	If not empty, airflow containers will attempt to install the additional requirements specified by this variable.

Was this entry helpful?



Want to be a part of Apache Airflow?

Join community



[Donate](#)[Thanks](#)[Security](#)

© The

Apache

Software  
Apache Airflow, Apache, Airflow, the Airflow logo, and the Apache feather logo are either

registered trademarks or trademarks of The Apache Software Foundation. All other products or

name brands are trademarks of their respective holders, including The Apache Software

Foundation.